

# Thomy Phan

## Curriculum Vitae

---

### Contact Information

Email thomy.phan@usc.edu  
Website thomyphan.github.io

---

### Research Interests

Multi-Agent Systems, Reinforcement Learning, Optimization

---

### Education

- 2018 – 2023 **Ph.D. in Computer Science**, *LMU Munich*, Germany
- Thesis: "*Emergence and Resilience in Multi-Agent Reinforcement Learning*"
  - Thesis committee: Claudia Linnhoff-Popien, Sven Koenig, Long Tran-Thanh
  - Based on work published at AAMAS, AAAI, IJCAI, NeurIPS, and ICML
- 2015 – 2017 **M.Sc. in Computer Science**, *LMU Munich*, Germany
- Focus on artificial intelligence, data science, and autonomous systems
  - Master thesis: "*EVADE: Emergent Value Function Approximation for Distributed Environments*"
  - Supervision: Claudia Linnhoff-Popien, Lenz Belzner
  - Results published at AAMAS 2018 as a **full conference paper**
- 2011 – 2015 **B.Sc. in Computer Science**, *Munich University of Applied Sciences*, Germany
- Collaborative study program (duales Studium) with the City of Munich
  - Focus on software development and image processing
  - Bachelor thesis: "*Quantification and Feature Extraction of 3D Single-Molecule Switching Microscopy Data*"
  - Supervision: Alfred Nischwitz, Joerg Bewersdorf
  - Practical work done at Bewersdorf Lab, Yale School of Medicine
  - Results published in Cell 2016 as a **journal paper (cover story)**

---

### Professional Experience

- 2023 – Present **Postdoctoral Scholar**, *University of Southern California*, Los Angeles, CA, USA
- Advisor: Sven Koenig
  - Focus on multi-agent learning for optimization.
- 2018 – 2023 **Research Assistant**, *LMU Munich*, Germany
- Advisor: Claudia Linnhoff-Popien
  - Focus on emergence and resilience in multi-agent systems.
- 2015 – 2018 **Software Developer**, *it@M, City of Munich*, Germany
- Development, maintainance, and integration of business applications (part-time).

---

### Internships

- 2015 **Visiting Scholar in Research (3 Months)**, *Yale University*, New Haven, CT, USA
- Advisor: Joerg Bewersdorf
  - Focus on data analysis and feature extraction of cellular structures in super-resolution microscopy data. Co-authored publication in Cell 2016 (cover story).
- 2012 – 2014 **Working Student (Collaborative Study Program)**, *City of Munich*, Germany
- Regular internships focusing on IT architecture, project management, and software development during the semester holidays.

---

## Honors and Awards

- 2023 **Outstanding Reviewer (Top 10%)**, *Conference on Neural Information Processing Systems (NeurIPS)*, New Orleans, LA, USA
- 2022 **Premier Paper of AAMAS 2022**, *Journal of the International Foundation for Autonomous Agents and Multi-Agent Systems (JAAMAS)*  
Fast-track invitation to submit an extended version of our AAMAS 2022 paper "*Emergent Cooperation from Mutual Acknowledgment Exchange*" (main author).
- 2022 **Highlight Paper at the Workshop on Ad Hoc Teamwork**, *International Joint Conference on Artificial Intelligence (IJCAI)*, Vienna, Austria  
Recognition of our AAMAS 2022 paper "*Emergent Cooperation from Mutual Acknowledgment Exchange*" (main author).
- 2022 **Outstanding Reviewer (Top 10%)**, *International Conference on Machine Learning (ICML)*, Baltimore, MD, USA
- 2021 **ICAART 2021 – Springer Selection**, *Lecture Notes in Artificial Intelligence*  
Fast-track invitation to submit an extended version of our ICAART 2021 paper "*SAT-MARL: Specification Aware Training in Multi-Agent Reinforcement Learning*" (co-author).
- 2019 **DAAD Travel Grant for AAMAS 2019**, *German Academic Exchange Service*, Montreal, Canada
- 2016 **Best Bachelor Award**, *Rohde & Schwarz GmbH & Co. KG*, Munich, Germany
- 2016 **Award for an Outstanding Bachelor Thesis in the Field of Image Processing**, *Stemmer Imaging GmbH*, Munich, Germany
- 2012 – 2017 **Scholarship**, *Studienstiftung des Deutschen Volkes*, Munich, Germany  
In Germany, the top 0.5% of university or high school students get selected for funding by the German Academic Scholarship Foundation.

---

## Research Projects

- 2024 – Present **Causal Foundations of Decision-Making and Learning**, *National Science Foundation*, Los Angeles, CA, USA  
Research on scalable causal reinforcement learning.
- 2023 – Present **AI4OPT – AI Institute for Advances in Optimization**, *National Science Foundation*, Los Angeles, CA, USA  
Research on multi-agent learning for optimization.
- 2023 **Dependability of Machine Learning in Industrial Robotics**, *Siemens AG*, Munich, Germany  
Research on robust machine learning in industrial robotics. Assisted acquisition.
- 2022 – 2024 **AI-Fusion – Evaluation of Emergence in Distributed Intelligent Systems**, *Bavarian Ministry of Economic Affairs, Regional Development, and Energy*, Munich, Germany  
Research on emergence in multi-agent learning in collaboration with Fraunhofer IKS. Assisted acquisition.
- 2022 **Validation and Verification of Modular Machine Learning Systems**, *Siemens AG*, Munich, Germany  
Research on modular machine learning. Assisted acquisition.
- 2020 – 2021 **Federated Learning in Industrial Environments**, *Siemens AG*, Munich, Germany  
Research on adaptive testing of federated learning systems. Assisted acquisition.

- 2019 **Dependability of Machine Learning in Industrial Environments**, *Siemens AG*, Munich, Germany  
Research on resilience in multi-agent reinforcement learning.
- 2018 **Coevolution in Machine Learning Based Industrial Environments**, *Siemens AG*, Munich, Germany  
Research on scenario coevolution in reinforcement learning.
- 2018 – 2023 **InnoMI – Innovation Center Mobile Internet**, *Bavarian Ministry of Economic Affairs, Regional Development, and Energy*, Munich, Germany  
Research on innovative mobile and distributed applications.

## Selected Publications

### Conferences

Extended abstracts ( $\leq 3$  pages) with a full conference or journal version are not listed.

- 2024 [C19] **Confidence-Based Curriculum Learning for Multi-Agent Path Finding**  
**Thomy Phan**, Joseph Driscoll, Justin Romberg, and Sven Koenig.  
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2024.  
To appear as a full paper.
- [C18] **Anytime Multi-Agent Path Finding Using Operator Parallelism in Large Neighborhood Search**  
Shao-Hung Chan, Zhe Chen, Dian-Lun Lin, Yue Zhang, Daniel Harabor, Sven Koenig, Tsung-Wei Huang, and **Thomy Phan**.  
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2024.  
To appear as an extended abstract.
- [C17] **Challenges for Reinforcement Learning in Quantum Computing**  
Philipp Altmann, Adelina Bärligea, Jonas Stein, Michael Kölle, Thomas Gabor, **Thomy Phan**, and Claudia Linnhoff-Popien.  
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2024.  
To appear as an extended abstract.
- [C16] **Adaptive Anytime Multi-Agent Path Finding Using Bandit-Based Large Neighborhood Search**  
**Thomy Phan**, Taoan Huang, Bistra Dilkina, and Sven Koenig.  
AAAI Conference on Artificial Intelligence (**AAAI**), 2024. To appear.
- 2023 [C15] **Attention-Based Recurrence for Multi-Agent Reinforcement Learning under Stochastic Partial Observability**  
**Thomy Phan**, Fabian Ritz, Philipp Altmann, Maximilian Zorn, Jonas Nüßlein, Michael Kölle, Thomas Gabor, and Claudia Linnhoff-Popien.  
International Conference on Machine Learning (**ICML**), pages 27840–27853, 2023.
- [C14] **CROP: Towards Distributional-Shift Robust Reinforcement Learning Using Compact Reshaped Observation Processing**  
Philipp Altmann, Leonard Feuchtinger, Fabian Ritz, Jonas Nüßlein, Claudia Linnhoff-Popien, and **Thomy Phan**.  
International Joint Conference on Artificial Intelligence (**IJCAI**), pages 3414–3422, 2023.
- 2022 [C13] **Emergent Cooperation from Mutual Acknowledgment Exchange**  
**Thomy Phan**, Felix Sommer, Philipp Altmann, Fabian Ritz, Lenz Belzner, and Claudia Linnhoff-Popien.  
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), pages 1047–1055, 2022.
- [C12] **Towards Anomaly Detection in Reinforcement Learning (Blue Sky Ideas)**  
Robert Müller, Steffen Illium, **Thomy Phan**, Tom Haider, and Claudia Linnhoff-Popien.  
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), pages 1799–1803, 2022.
- 2021 [C11] **VAST: Value Function Factorization with Variable Agent Sub-Teams**  
**Thomy Phan**, Fabian Ritz, Lenz Belzner, Philipp Altmann, Thomas Gabor, and Claudia Linnhoff-Popien.  
Advances in Neural Information Processing Systems (**NeurIPS**), pages 24018–24032, 2021.

- [C10] **Resilient Multi-Agent Reinforcement Learning with Adversarial Value Decomposition**  
**Thomy Phan**, Lenz Belzner, Thomas Gabor, Andreas Sedlmeier, Fabian Ritz, and Claudia Linnhoff-Popien.  
 AAAI Conference on Artificial Intelligence (AAAI), pages 11308–11316, 2021.
- [C9] **SAT-MARL: Specification Aware Training in Multi-Agent Reinforcement Learning**  
 Fabian Ritz, **Thomy Phan**, Robert Müller, Thomas Gabor, Andreas Sedlmeier, Marc Zeller, Jan Wieghardt, Reiner Schmid, Horst Sauer, Cornel Klein, and Claudia Linnhoff-Popien.  
 International Conference on Agents and Artificial Intelligence (ICAART), pages 28–37, 2021.
- 2020 [C8] **A Quantum Annealing Algorithm for Finding Pure Nash Equilibria in Graphical Games**  
 Christoph Roch, **Thomy Phan**, Sebastian Feld, Robert Müller, Thomas Gabor, Carsten Hahn, and Claudia Linnhoff-Popien.  
 International Conference on Computational Science (ICCS), pages 488–501, 2020.
- [C7] **Learning and Testing Resilience in Cooperative Multi-Agent Systems**  
**Thomy Phan**, Thomas Gabor, Andreas Sedlmeier, Fabian Ritz, Bernhard Kempter, Cornel Klein, Horst Sauer, Reiner Schmid, Jan Wieghardt, Marc Zeller, and Claudia Linnhoff-Popien.  
 International Conference on Autonomous Agents and Multiagent Systems (AAMAS), pages 1055–1063, 2020.
- 2019 [C6] **Adaptive Thompson Sampling Stacks for Memory Bounded Open-Loop Planning**  
**Thomy Phan**, Thomas Gabor, Robert Müller, Christoph Roch, and Claudia Linnhoff-Popien.  
 International Joint Conference on Artificial Intelligence (IJCAI), pages 5607–5613, 2019.
- [C5] **Subgoal-Based Temporal Abstraction in Monte-Carlo Tree Search**  
 Thomas Gabor, Jan Peter, **Thomy Phan**, Christian Meyer, and Claudia Linnhoff-Popien.  
 International Joint Conference on Artificial Intelligence (IJCAI), pages 5562–5568, 2019.
- [C4] **Scenario Co-Evolution for Reinforcement Learning on a Grid World Smart Factory Domain**  
 Thomas Gabor, Andreas Sedlmeier, Marie Kiermeier, **Thomy Phan**, Marcel Henrich, Monika Pichlmair, Bernhard Kempter, Cornel Klein, Horst Sauer, Reiner Schmid, and Jan Wieghardt.  
 Genetic and Evolutionary Computation Conference (GECCO), pages 898–906, 2019.
- [C3] **Distributed Policy Iteration for Scalable Approximation of Cooperative Multi-Agent Policies (Extended Abstract)**  
**Thomy Phan**, Kyrill Schmid, Lenz Belzner, Thomas Gabor, Sebastian Feld, and Claudia Linnhoff-Popien.  
 International Conference on Autonomous Agents and Multiagent Systems (AAMAS), pages 2162–2164 (extended abstract), 2019.
- [C2] **Memory Bounded Open-Loop Planning in Large POMDPs Using Thompson Sampling**  
**Thomy Phan**, Lenz Belzner, Marie Kiermeier, Markus Friedrich, Kyrill Schmid, and Claudia Linnhoff-Popien.  
 AAAI Conference on Artificial Intelligence (AAAI), pages 7941–7948, 2019.
- 2018 [C1] **Leveraging Statistical Multi-Agent Online Planning with Emergent Value Function Approximation**  
**Thomy Phan**, Lenz Belzner, Thomas Gabor, and Kyrill Schmid.  
 International Conference on Autonomous Agents and Multiagent Systems (AAMAS), pages 730–738, 2018.

## Journals

- 2024 [J5] **Emergent Cooperation from Mutual Acknowledgment Exchange in Multi-Agent Reinforcement Learning**  
**Thomy Phan**, Felix Sommer, Fabian Ritz, Philipp Altmann, Jonas Nüßlein, Michael Kölle, Lenz Belzner, and Claudia Linnhoff-Popien.  
 Journal on Autonomous Agents and Multi-Agent Systems (JAAMAS), 2024. Invited from AAMAS 2022.

- 2022 [J4] **Specification Aware Multi-Agent Reinforcement Learning**  
 Fabian Ritz, **Thomy Phan**, Robert Müller, Thomas Gabor, Andreas Sedlmeier, Marc Zeller, Jan Wieghardt, Reiner Schmid, Horst Sauer, Cornel Klein, and Claudia Linnhoff-Popien.  
*Springer Book of ICAART 2021*, pages 3–21, 2022. Invited from ICAART 2021.
- 2021 [J3] **Productive Fitness in Diversity-Aware Evolutionary Algorithms**  
 Thomas Gabor, **Thomy Phan**, and Claudia Linnhoff-Popien.  
*Natural Computing*, 20(3): 363–376, 2021.
- 2020 [J2] **The Scenario Coevolution Paradigm: Adaptive Quality Assurance for Adaptive Systems**  
 Thomas Gabor, Andreas Sedlmeier, **Thomy Phan**, Fabian Ritz, Marie Kiermeier, Lenz Belzner, Bernhard Kempter, Cornel Klein, Horst Sauer, Reiner Schmid, Jan Wieghardt, Marc Zeller, and Claudia Linnhoff-Popien.  
*International Journal on Software Tools for Technology Transfer (STTT)*, 22(4): 457–476, 2020.
- 2016 [J1] **Ultra-High Resolution 3D Imaging of Whole Cells (Cover Story)**  
 Fang Huang, George Sirinakis, Edward S Allgeyer, Lena K Schroeder, Whitney C Duim, Emil B Kromann, **Thomy Phan**, Felix E Rivera-Molina, Jordan R Myers, Irnov Irnov, Mark Lessard, Yongdeng Zhang, Mary Ann Handel, Christine Jacobs-Wagner, C Patrick Lusk, James E Rothman, Derek Toomre, Martin J Booth, and Joerg Bewersdorf.  
*Cell*, 166(4): 1028–1040, 2016.

### Workshops

Workshop papers with a conference or journal version are not listed.

- 2023 [W3] **DIRECT: Learning from Sparse and Shifting Rewards Using Discriminative Reward Co-Training**  
 Philipp Altmann, **Thomy Phan**, Fabian Ritz, Thomas Gabor, and Claudia Linnhoff-Popien.  
 AAMAS Workshop Adaptive and Learning Agents (**ALA**), 2023.
- 2020 [W2] **The Holy Grail of Quantum Artificial Intelligence: Major Challenges in Accelerating the Machine Learning Pipeline**  
 Thomas Gabor, Leo Sünnkel, Fabian Ritz, **Thomy Phan**, Lenz Belzner, Christoph Roch, Sebastian Feld, and Claudia Linnhoff-Popien.  
 ICSE Workshop on Quantum Software Engineering (**Q-SE**), pages 456–461, 2020.
- [W1] **A Distributed Policy Iteration Scheme for Cooperative Multi-Agent Policy Approximation**  
**Thomy Phan**, Lenz Belzner, Kyrill Schmid, Thomas Gabor, Fabian Ritz, Sebastian Feld, and Claudia Linnhoff-Popien.  
 AAMAS Workshop Adaptive and Learning Agents (**ALA**), 2020.

---

## Academic Activities

### Organizing Committee

- 2019 International Symposium on Applied Artificial Intelligence (ISAAI)

### Program Committee

- 2023 – 2024 International Joint Conference on Artificial Intelligence (IJCAI – Main Track)
- 2021 – 2024 AAAI Conference on Artificial Intelligence (AAAI)
- 2023 International Joint Conference on Artificial Intelligence (IJCAI – AI and Social Good)
- 2023 European Conference on Artificial Intelligence (ECAI)
- 2023 International Conference on Autonomous Agents and Multiagent Systems (AAMAS – Blue Sky Ideas)

### Reviewer

- 2022 – 2024 International Conference on Machine Learning (ICML)
- 2024 International Conference on Automated Planning and Scheduling (ICAPS)
- 2024 International Conference on Learning Representations (ICLR)

- 2024 AAAI 2024 Workshop on Cooperative Multi-Agent Systems Decision-making and Learning (CMASDL)
- 2022 – 2023 Conference on Neural Information Processing Systems (NeurIPS)
- 2023 International Conference on Autonomous Agents and Multiagent Systems (AAMAS – Main Track)
- 2018, 2022 International Symposium On Leveraging Applications of Formal Methods (ISoLA)
- 2021 PLOS ONE Journal
- 2020 International Journal on Software Tools for Technology Transfer (STTT)

---

## Teaching

- Spring 2024 **CSCI 599: Autonomous Decision-Making**, *University of Southern California*, Los Angeles, CA, USA
  - Special topics lecture on planning, reinforcement learning, and multi-agent systems
  - Primary instructor
  - Syllabus and registration count provided at <https://classes.usc.edu/term-20241/course/csci-599/>
- 2019 – 2023 **Autonomous Systems**, *LMU Munich*, Germany
  - Practical course for 12 – 18 master students on planning and reinforcement learning
  - Primary supervising assistant until summer semester 2022
  - Syllabus and registration count for summer semester 2022 provided at <https://uni2work.ifi.lmu.de/course/S22/IfI/ASP>
- 2019 – 2023 **Working Group "Artificial Intelligence"**, *LMU Munich*, Germany
  - Voluntary working group for more than 100 bachelor and master students on current AI topics
  - Primary supervising assistant
  - Syllabus and registration count for summer semester 2022 provided at <https://uni2work.ifi.lmu.de/course/S22/IfI/AIAG>
- 2018 – 2019 **Mobile and Distributed Systems**, *LMU Munich*, Germany
  - Practical course for 12 – 18 master students on mobile app development and on-device machine learning
  - Secondary supervising assistant
- 2018 – 2023 **Student Mentoring**, *LMU Munich*, Germany
  - 26 master theses
  - 24 bachelor theses
  - 8 individual research projects

---

## Talks

### Invited Talks

- 07/2022 **Emergent Cooperation from Mutual Acknowledgment Exchange**, Workshop on Ad Hoc Teamwork at IJCAI 2022 (virtual)  
Highlight presentation of our AAMAS 2022 paper "*Emergent Cooperation from Mutual Acknowledgment Exchange*" (main author).
- 06/2021 **Stability in AI-Systems**, Digitale Stadt München e.V., Germany (virtual)  
DigiTalk event on Safe Intelligence of the Digital City Association of Munich.
- 12/2020 **Artificial Intelligence – How Do Robots Learn?**, Gymnasium Berchtesgaden, Germany (virtual)  
P-seminar talk for high school students and the Junior Science Café.
- 03/2019 **Building Autonomous Systems with AI**, *University of Augsburg*, Germany  
AI workshop for students of the Software Engineering Elite Graduate Program.

### [Presentations at Conferences as Main Author](#)

- 07/2023 **Attention-Based Recurrence for Multi-Agent Reinforcement Learning under Stochastic Partial Observability**, Honolulu, Hawaii, USA  
International Conference on Machine Learning (ICML).
- 05/2022 **Emergent Cooperation from Mutual Acknowledgment Exchange**, virtual  
International Conference on Autonomous Agents and Multiagent Systems (AAMAS).
- 12/2021 **VAST: Value Function Factorization with Variable Agent Sub-Teams**, virtual  
Conference on Neural Information Processing Systems (NeurIPS).
- 02/2021 **Resilient Multi-Agent Reinforcement Learning with Adversarial Value Decomposition**, virtual  
AAAI Conference on Artificial Intelligence (AAAI).
- 05/2020 **Learning and Testing Resilience in Cooperative Multi-Agent Systems**, virtual  
International Conference on Autonomous Agents and Multiagent Systems (AAMAS).
- 08/2019 **Adaptive Thompson Sampling Stacks for Memory Bounded Open-Loop Planning**, Macao, China  
International Joint Conference on Artificial Intelligence (IJCAI).
- 05/2019 **Distributed Policy Iteration for Scalable Approximation of Cooperative Multi-Agent Policies**, Montreal, Canada  
International Conference on Autonomous Agents and Multiagent Systems (AAMAS).
- 02/2019 **Memory Bounded Open-Loop Planning in Large POMDPs using Thompson Sampling**, Honolulu, Hawaii, USA  
AAAI Conference on Artificial Intelligence (AAAI).
- 07/2018 **Leveraging Statistical Multi-Agent Online Planning with Emergent Value Function Approximation**, Stockholm, Sweden  
International Conference on Autonomous Agents and Multiagent Systems (AAMAS).

### [Other Talks](#)

- 2019 **Unitag: Artificial Intelligence – How Do Robots Learn?**, *LMU Munich*, Germany  
University event for gifted high school students from Upper Bavaria (Oberbayern).